

# Introduction to Artificial Intelligence

Daniel Deng

## 1 Search Problems

**Definition 1.1** (Reflex Agent). A reflex agent chooses actions based on its current perception of the world.

**Definition 1.2** (Planning Agent). A planning agent chooses actions based on hypothesized consequences of actions.

**Definition 1.3** (Search Problem). A search problem consists of a state space, a successor function, a start state, and a goal test.

## 2 Search Algorithms

### 2.1 Heuristics

**Definition 2.1** (Heuristic). A heuristic  $h(n)$  is a function that estimates the distance from state  $n$  to the goal state for a particular search problem. It is often solutions of relaxed problems.

**Definition 2.2** (Admissibility). A heuristic is admissible, or optimistic, if  $0 \leq h(n) \leq h^*(n)$  where  $h^*$  is the true cost to goal state.

**Definition 2.3** (Consistency). A heuristic is consistent if  $h(n) - h(n+1) \leq c(n, n+1)$  where  $c$  is the cost between states  $n$  and  $n+1$ .

*Remark.* Consistency necessarily implies admissibility.

Table 1: Search algorithms.

	Fringe	Complete	Optimal	Time	Space
Depth-First Search	Stack	<i>iff</i> no cycle	No	$O(b^m)$	$O(bm)$
Breadth-First Search	Queue	Yes	<i>iff</i> uniform cost	$O(b^s)^1$	$O(b^s)^1$
Uniform Cost Search	PQ $(g(n))^2$	<i>iff</i> positive cost	Yes	$O(b^{c^*/\epsilon})^3$	$O(b^{c^*/\epsilon})^3$
Greedy Search	PQ $(h(n))$	-	No	-	-
$A^*$ Tree Search	PQ $(h(n) + g(n))$	-	<i>iff</i> $h(n)$ admissible	-	-
$A^*$ Graph Search <sup>4</sup>	PQ $(h(n) + g(n))$	-	<i>iff</i> $h(n)$ consistent	-	-

<sup>1</sup>  $s$  = depth of solution.<sup>2</sup>  $g(n)$  = cumulative path cost.<sup>3</sup>  $c^*/\epsilon$  = effective solution depth ( $c^*$  = cost of the cheapest solution;  $\epsilon$  = minimum cost of cost-contour arcs).<sup>4</sup> Compared to tree search, graph search keeps a closed set of expanded states to check against to prevent duplicate expansions.

*Remark.* Implementation of search algorithms differ only in fringe strategies.

### 3 Constrained Satisfaction Problems

**Definition 3.1** (Constrained Satisfaction Problems). Constrained Satisfaction Problems (CSPs) are a type of **identification problem** defined by variable  $X_0, \dots, X_n$  with values from a domain  $D$  that satisfies a set of constraints.

---

**Algorithm 1:** Backtracking search.

---

**Input:** A constraint satisfaction problem  $P$ .

**Output:** A complete assignment  $A$ .

```

1 Function BS( $P$ ):
2   return EXPLORE ( $P$ , {})
3 Function EXPLORE( $P$ ,  $A$ ):
4   if  $A$  is complete then
5     return  $A$ 
6    $unassigned \leftarrow$  an unassigned VARIABLES( $P$ )
7   foreach  $value \in$  DOMAIN( $unassigned$ ) do
8     if  $value$  is consistent with all CONSTRAINTS( $P$ ) then
9       add { $unassigned \leftarrow value$ } to  $A$ 
10       $attempt \leftarrow$  EXPLORE( $P$ ,  $A$ )
11      if  $attempt$  failed then
12        remove { $unassigned \leftarrow value$ } from  $A$ 
13      else
14        return  $attempt$ 
15   return failed

```

---

### 3.1 Filtering

**Definition 3.2** (Arc Consistency).

Arc  $X \rightarrow Y$  is consistent  $\Leftrightarrow$   
 $(\forall x \in D_x)(\exists y \in D_y)(y \text{ can be assigned to } Y \text{ without violating a constraint.})$

---

**Algorithm 2:** Arc consistency filtering.

---

**Input:** A constraint satisfaction problem  $P$ .

```

1 Function AC-3( $P$ ):
2    $Q \leftarrow$  empty Queue
3   enqueue all arcs  $\in P$ 
4   while  $Q$  is not empty do
5      $(X_i, X_j) \leftarrow$  dequeue from  $Q$ 
6     if FILTER( $X_i, X_j$ ) is successful then
7       foreach  $X_k \in \text{NEIGHBORS}(X_i)$  do
8         enqueue  $(X_k, X_i)$ 
9 Function FILTER ( $tail, head$ ):
10   $result \leftarrow$  false
11  foreach  $value \in \text{DOMAIN}(tail)$  do
12    if  $value$  violates some constraint with all values in  $\text{DOMAIN}(head)$  then
13      delete  $value$  from  $\text{DOMAIN}(tail)$ 
14       $result \leftarrow$  true
15  return  $result$ 

```

---

### 3.2 Ordering

**Definition 3.3** (Minimum Remaining Values). The MRV policy chooses an unassigned variable that has the fewest valid remaining values in order to induce backtracking earlier and reduce potential node expansions.

**Definition 3.4** (Least Constraining Value). The LCV policy chooses a value assignment that violates the least amount of constraints, which requires additional computation such as running arc consistency test on each value.

### 3.3 Structure

Given a tree-structured CSP, represent it as a directed acyclic graph. Enforcing arc consistency in reverse topological order then assigning in topological order ensures a runtime

of  $O(nd^2)$  (as opposed to  $O(d^n)$  in the general case).  
 TODO: nearly tree-like CSPs and tree decomposition.

## 4 Local Search

Improve a single option until no further improvements can be made.

*Remark.* Generally, local search is faster and more memory efficient at the expense of completeness and optimality.

### 4.1 Iterative Algorithm for CSP/Hill Climbing

Starting with a "complete" state, randomly select any conflicted variables and reassign values using min-conflicts heuristics.

*Remark.* Efficiency of the algorithm depends on  $R = \frac{\text{number of constraints}}{\text{number of variables}}$ ; computation time is approximately constant time except when  $R$  approaches the *critical ratio*.

### 4.2 Simulated Annealing

---

**Algorithm 3:** Simulated annealing.

---

**Input:** A problem  $P$  and a schedule/mapping from time to "temperature"  $T$ .

**Output:** A solution state.

*/\* Escape local maxima by allowing downhill movement based on a "temperature"-dependent probabilistic function. \*/*

```

1 current ← initial state of  $P$ 
2 for  $t \leftarrow 1$  to  $\infty$  do
3    $temp \leftarrow T[t]$ 
4   if  $temp = 0$  then
5     return current
6   else
7      $next \leftarrow$  a randomly selected successor of current
8      $\Delta \leftarrow \text{VALUE}[next] - \text{VALUE}[current]$ 
9     if  $\Delta > 0$  then
10       $current \leftarrow next$ 
11    else
12       $currnet \leftarrow next$  with probability  $e^{\frac{\Delta}{temp}}$ 
```

---

## 5 Genetic Algorithms

Keep best  $N$  hypotheses at each step (selection) based on a fitness function and have pairwise crossover operations (and, optionally, mutation operations) to generate a new set of hypotheses.

## 6 Games

---

**Algorithm 4:** Minimax with alpha-beta pruning.

---

**Input:** A game state  $S$   
**Output:** The root minimax value.  
 /\* initialize  $\alpha \leftarrow -\infty$  and  $\beta \leftarrow \infty$  \*/

```

1 Function VALUE( $S, \alpha, \beta$ ):
2   if  $S$  is a terminal state then
3     return known terminal value
4   if the agent is maximizing then
5     return MAX-VALUE( $S, \alpha, \beta$ )
6   if the agent is minimizing then
7     return MIN-VALUE( $S, \alpha, \beta$ )

8 Function MAX-VALUE( $S, \alpha, \beta$ ):
9    $v \leftarrow -\infty$ 
10  foreach successor  $S'$  of  $S$  do
11     $v \leftarrow \text{MAX}(v, \text{VALUE}(S'))$ 
12    if  $v \geq \beta$  then
13      return  $v$ 
14     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
15  return  $v$ 

16 Function MIN-VALUE( $S, \alpha, \beta$ ):
17   $v \leftarrow \infty$ 
18  foreach successor  $S'$  of  $S$  do
19     $v \leftarrow \text{MIN}(v, \text{VALUE}(S'))$ 
20    if  $v \leq \alpha$  then
21      return  $v$ 
22     $\beta \leftarrow \text{MIN}(\beta, v)$ 
23  return  $v$ 

```

---

---

**Algorithm 5:** Expectimax.

---

**Input:** A game state  $S$

**Output:** The root minimax value.

```

1 Function VALUE( $S$ ):
2   if  $S$  is a terminal state then
3     return known terminal value
4   if the agent is maximizing then
5     return MAX-VALUE( $S$ )
6   if the agent is randomizing then
7     return EXP-VALUE( $S$ )

8 Function MAX-VALUE( $S$ ):
9    $v \leftarrow -\infty$ 
10  foreach successor  $S'$  of  $S$  do
11     $v \leftarrow \text{MAX}(v, \text{VALUE}(S'))$ 
12  return  $v$ 

13 Function EXP-VALUE( $S$ ):
14   $v \leftarrow 0$ 
15  foreach successor  $S'$  of  $S$  do
16     $v \leftarrow \mathbb{E}[\text{VALUE}(S')] // \text{expected value}$ 
17  return  $v$ 

```

---

*Remark.* Typically, performing Minimax and Expectimax to terminal states are too costly, which can be solved by terminating early and estimating the state utility with evaluation functions.

## 7 Markov Decision Processes

*Remark.* Finite horizons (finite timestep before an agent terminates) and/or discount factors ( $\gamma$ ) ensure an agent terminates in MDP.

**Definition 7.1** (Transition Function).  $T(s, a, s') = P(s' \mid s, a)$

**Definition 7.2** (Value Iteration). Initialize  $V_0(s) \leftarrow 0$  for all  $s \in S$ . Compute  $\forall s \in S$

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')] \quad (\text{value iteration})$$

until  $V_i(s)$  converges for all  $s \in S$  (yields  $V^*$ ). Then, compute  $\forall s \in S$

$$\pi^*(s) \leftarrow \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')] \quad (\text{policy extraction})$$

**Definition 7.3** (Policy Iteration). Define an initial policy  $\pi_0$  (can be arbitrary, but ideality close to the optimal policy). Then, iteratively solve  $\forall s \in S$

$$V^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V^{\pi_i}(s')] \quad (\text{policy evaluation})$$

$$\pi_{i+1}(s) \leftarrow \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')] \quad (\text{policy improvement})$$

until  $\pi(s)$  converges for all  $s \in S$  (yields  $\pi^*$ ).

*Remark.* Policy evaluation solves a system of  $|S|$  linear equations.

*Remark.* Policy iteration converges faster than value iteration.