Introduction to Artificial Intelligence

Daniel Deng

1   Agents

Definition 1.1 (Reflex Agent). A reflex agent chooses actions based on its current perception of the world.

Definition 1.2 (Planning Agent). A planning agent chooses actions based on hypothesized consequences of actions.

2 General Search Problems

2.1 Heuristics

Definition 2.1 (Heuristic)**.** A heuristic $h(n)$ is a function that estimates the distance from state n to the goal state for a particular search problem. It is often solutions of relaxed problems.

1. A heuristic is <u>admissible</u> if $0 \leq h(n) \leq h^*(n)$ where h^* is the true cost to goal state;

2. A heuristic is <u>consistent</u> if $h(n) - h(n+1) \leq c(n, n+1)$ where c is the cost between states n and $n+1$.

Remark. Consistency necessarily implies admissibility.

2.2 Search Algorithms

Table 1: Search algorithms.

	Fringe	Complete	Optimal	Time	Space
Depth-First Search	Stack	*iff* no cycle	No	$O(b^m)$	$O(bm)$
Breadth-First Search	Queue	Yes	*iff* uniform cost	$O(b^s)$[1]	$O(b^s)$[1]
Uniform Cost Search	PQ $(g(n))$[2]	*iff* positive cost	Yes	$O(b^{c^*/\epsilon})$[3]	$O(b^{c^*/\epsilon})$[3]
Greedy Search	PQ $(h(n))$	-	No	-	-
A^* Tree Search	PQ $(h(n) + g(n))$	-	*iff* $h(n)$ admissible	-	-
A^* Graph Search[4]	PQ $(h(n) + g(n))$	-	*iff* $h(n)$ consistent	-	-

[1] s = depth of solution.
[2] $g(n)$ = cumulative path cost.
[3] c^*/ϵ = effective solution depth (c^* = cost of the cheapest solution; ϵ = minimum cost of cost-contour arcs).
[4] Compared to tree search, graph search keeps a closed set of expanded states to check against to prevent duplicate expansions.

Remark. Implementation of search algorithms differ only in fringe strategies.

3 Constrained Satisfaction Problems

Definition 3.1 (Constrained Satisfaction Problems)**.** Constrained Satisfaction Problems (CSPs) are a type of **identification problem** defined by variable $X_0, \ldots, X_n$ with values from a domain D that satisfies a set of constrains.

Algorithm 1: Backtracking search.

    Input: A constraint satisfaction problem P.
    Output: A complete assignment A.

1 Function BS(P)**:**
2   **return** EXPLORE *(P, {})*

3 Function EXPLORE(P, A)**:**
4   **if** *A is complete* **then**
5   **return** A

6   *unassigned* $\leftarrow$ an unassigned VARIABLES(P)
7   **foreach** *value* $\in$ DOMAIN(*unassigned*) **do**
8   **if** *value is consistent with all* CONSTRAINTS(P) **then**
9   add $\{unassigned \leftarrow value\}$ to A
10   *attempt* $\leftarrow$ EXPLORE(P, A)
11   **if** *attemp failed* **then**
12   remove $\{unassigned \leftarrow value\}$ from A
13   **else**
14   **return** *attempt*

15   **return** *failed*

3.1 Filtering

Definition 3.2 (Arc Consistency)**.**

Arc $X \to Y$ is consistent $\Leftrightarrow$

$(\forall x \in D_x)(\exists y \in D_y)(y$ can be assigned to Y without violating a constraint.$)$

Algorithm 2: Arc consistency filtering.

Input: A constraint satisfaction problem P.

/* $O(|S||A|^3)$ runtime   */

1 **Function** AC-3(P):
2      $Q \leftarrow$ empty Queue
3   enqueue all arcs $\in P$
4   **while** Q *is not empty* **do**
5          $(X_i, X_j) \leftarrow$ dequeue from Q
6   **if** FILTER(X_i, X_j) *is successful* **then**
7   **foreach** $X_k \in$ NEIGHBORS(X_i) **do**
8   enqueue (X_k, X_i)

9 **Function** FILTER *(tail, head)*:
10      $result \leftarrow$ false
11   **foreach** $value \in$ DOMAIN($tail$) **do**
12   **if** *value violates some constraint with all values in* DOMAIN($head$) **then**
13   delete $value$ from DOMAIN($tail$)
14              $result \leftarrow$ true
15   **return** $reuslt$

3.2 Ordering

Definition 3.3 (Minimum Remaining Values)**.** The MRV policy chooses an unassigned variable that has the fewest valid remaining values in order to induce backtracking earlier and reduce potential node expansions.

Definition 3.4 (Least Constraining Value)**.** The LCV policy chooses a value assignment that violates the least amount of constraints, which requires additional computation such as running arc consistency test on each value.

3.3 Structure

Given a tree-structured CSP, represent it as a directed acyclic graph. Enforcing arc consistency in reverse topological order then assigning in topological order ensures a runtime of $O(nd^2)$ (as opposed to $O(d^n)$ in the general case).
TODO: nearly tree-like CSPs and tree decomposition.

4   Local Search

Definition 4.1 (Local Search). A search strategy that improve a single option until no further improvements can be made. Typically, local search is faster and more memory efficient than other search algorithms but is generally neither complete nor optimal.

Definition 4.2 (Hill Climbing). A CSP strategy that randomly selects a conflicting variable and reassign values using min-conflicts heuristics.

Remark. Efficiency of the algorithm depends on $R = \frac{\text{number of constraints}}{\text{number of variables}}$; computation time is approximately constant time except when R approaches the *critical ratio*.

4.1   Simulated Annealing

Algorithm 3: Simulated annealing.

   Input: A problem P and a schedule/mapping from time to "temperature" T.
   Output: A solution state.
   /* Escape local maxima by allowing downhill movement based on a
     "temperature"-dependent probabilistic function.   */

1   $current \leftarrow$ initial state of P
2   **for** $t \leftarrow 1$ *to* ∞ **do**
3      $temp \leftarrow T[t]$
4   **if** $temp = 0$ **then**
5   **return** $current$
6   **else**
7         $next \leftarrow$ a randomly selected successor of $current$
8         $\Delta \leftarrow \text{VALUE}[next]$ - $\text{VALUE}[current]$
9   **if** $\Delta > 0$ **then**
10           $current \leftarrow next$
11   **else**
12           $currnet \leftarrow next$ with probability $e^{\frac{\Delta}{temp}}$

4.2   Genetic Algorithm

Definition 4.3 (Genetic Algorithm). A strategy that keeps the best N hypothesis at each step based on a fitness function and generate next generation using pairwise cross-over operations (and, optionally, mutation operations).

5 Games

Definition 5.1 (Games)**.** Games are multi-agent search problems that could be zero-sum (adversarial) or general sum.

5.1 Minimax

Definition 5.2 (Minimax)**.** A zero-sum-game algorithm that assumes the opponent is an optimal adversary.

Algorithm 4: Minimax with alpha-beta pruning.

    Input: A game state S
    Output: The root minimax value.
    /* initialize $\alpha \leftarrow -\infty$ and $\beta \leftarrow \infty$ */

1 **Function** VALUE(S, α, β):
2   **if** *S is a terminal state* **then**
3   **return** *knwon temrinal value*

4   **if** *the agent is maximizing* **then**
5   **return** MAX-VALUE(S, α, β)

6   **if** *the agent is minimizing* **then**
7   **return** MIN-VALUE *(S, α, β)*

8 **Function** MAX-VALUE(S, α, β):
9 $v \leftarrow -\infty$
10   **foreach** *successor S' of S* **do**
11 $v \leftarrow$ MAX(v, VALUE(S'))
12   **if** $v \geq \beta$ **then**
13   **return** v
14 $\alpha \leftarrow$ MAX(α, v)
15   **return** v

16 **Function** MIN-VALUE(S, α, β):
17 $v \leftarrow \infty$
18   **foreach** *successor S' of S* **do**
19 $v \leftarrow$ MIN(v, VALUE(S'))
20   **if** $v \leq \alpha$ **then**
21   **return** v
22 $\beta \leftarrow$ MIN(β, v)
23   **return** v

5.2 Expectimax

Definition 5.3 (Expectimax)**.** A zero-sum-game algorithm that assumes the opponent acts based on some probability distribution.

Algorithm 5: Expectimax.

Input: A game state S
Output: The root minimax value.

1 **Function** VALUE(S):
2   **if** *S is a terminal state* **then**
3   **return** *knwon temrinal value*
4   **if** *the agent is maximizing* **then**
5   **return** MAX-VALUE(S)
6   **if** *the agent is randomizing* **then**
7   **return** EXP-VALUE(S)

8 **Function** MAX-VALUE(S):
9     $v \leftarrow -\infty$
10   **foreach** *successor S' of S* **do**
11        $v \leftarrow$ MAX(v, VALUE(S'))
12   **return** v

13 **Function** EXP-VALUE(S):
14     $v \leftarrow 0$
15   **foreach** *successor S' of S* **do**
16        $v \leftarrow \mathbb{E}\left[\text{VALUE}(S')\right]$ // expected value
17   **return** v

---

6 Markov Decision Processes

Remark. Finite horizons (finite timestep before an agent terminates) and/or discount factors (γ) ensure an agent terminates in MDP.

Definition 6.1 (Transition Function). $T(s, a, s') = P(s' \mid s, a)$

6.1 Value Iteration

Algorithm 6: Value Iteration

    Input: A MDP (S, A, R, T, γ).
    Output: The optimal policy $\pi^*(s)$ for all $s \in S$.

1 Function POLICY-EXTRACTION():
2 $V^* \leftarrow$ VALUE-ITERATION() // optimal value
3   **foreach** $s \in S$ **do**
4 $\pi^*(s) \leftarrow \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V_k(s')]$
5   **return** $\pi^*(s) \mid \forall s \in S$

6 Function VALUE-ITERATION():
    // $O(|S|^2|A|)$ runtime
7   Initialize $V_0(s) \leftarrow 0$ for all $s \in S$
8   **while** $V_{k+1}(s) \neq V_k(s) \mid \forall s \in S$ **do**
       // repeat until values converge
9 $V_{k+1}(s) \leftarrow \underset{a}{\max} \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V_k(s')]$
10   **return** $V^*(s) \leftarrow V_{k+1}(s) \mid \forall s \in S$

6.2 Policy Iteration

Definition 6.2 (Policy Iteration). Define an initial policy π_0 (can be arbitrary, but ideality close to the optimal policy). Then, iteratively solve $\forall s \in S$

$$V^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s')[R(s, \pi_i(s), s') + \gamma V^{\pi_i}(s')] \qquad \text{(policy evaluation)}$$

$$\pi_{i+1}(s) \leftarrow \underset{a}{\operatorname{argmax}} \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V^{\pi_i}(s')] \qquad \text{(policy improvement)}$$

until $\pi(s)$ converges for all $s \in S$ (yields pi^*.

Remark. Policy evaluation solves a system of $|S|$ liear equations.

Remark. Policy iteration converges faster than value iteration.

7 Reinforcement Learning

Remark. Reinforcement learning operates on MDP problems where the T and R functions are unknown.

7.1 Model-Based Learning

Definition 7.1 (Model-Based RL)**.** An algorithm that counts and normalizes sample outcomes s' for each s, a to construct $\hat{T}(s, a, s')$ and discovers each $\hat{R}(s, a, s')$ through exploration. The approximated MDP is then solved by value or policy iteration.

7.2 Model-Free Learning

Definition 7.2 (Direct Evaluation)**.** An algorithm that fixes some policy π and empirically computes $V^\pi(s)$ for all $s \in S$ by averaging total sample utility for a given state.

Remark. Direct evaluation wastes information about state connections and will take a long time to learn.

Definition 7.3 (Temporal Difference Learning)**.**

$$V^\pi(s) \leftarrow (1 - \alpha)V^\pi(s) + \alpha[R(s, \pi(s), s') + \gamma V^\pi(s')]$$

Remark. Can't extract policy.

Definition 7.4 (Q-Learning)**.**

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha[R(s, a, s') + \gamma \max_{a'} Q(s', a')]$$

Remark. Q-learning is an example of off-policy learning, in which the algorithm can learn the optimal policy even by taking sub-optimal or random actions.

Definition 7.5 (Approximate Q-Learning)**.** Represent q-values as weighted sums of features $Q(s, a) := \vec{w} \cdot \vec{f}(s, a)$. The update rule for Q-Learning then becomes

$$\Delta \leftarrow [R'(s, a, s') + \gamma \max_{a'} Q(s', a')] - Q(s, a)$$

$$w_i \leftarrow w_i + \alpha \cdot \Delta \cdot f_i(s, a)$$

Definition 7.6 (ϵ-Greedy Policies)**.** Define some probability $0 \leq \epsilon \leq 1$ to act randomly and explore. ϵ should be lowered over time to favor more exploitation as the learning becomes complete.

Definition 7.7 (Exploration Function)**.**

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha[R(s, a, s') + \gamma \max_{a'} f(s', a')]$$

$$f(s, a) = Q(s, a) + \frac{\text{const.}}{\text{count}(Q(s,a))}$$

Remark. $f(s, a)$ shown here is only a common example where the "bonus" for exploration diminishes as q-states are explored more and more.

8 Bayesian Network

Definition 8.1 (Conditional Independence). $\forall x, y, z : P(x|z, y) = P(x|z)$ (i.e., $x \perp\!\!\!\perp y \mid z$)

Definition 8.2 (Bayes' Net). Bayes' Net is a graphic model (DAG) that describes complex joint distributions using simple, local distributions (conditional probabilities)

- Nodes—variables (with domains) [assigned = observed, unassigned = unobserved]

- Arcs—interactions (i.e., encode conditional independence when lacking arrows) $\implies$
  $P(x_i|x_1, \ldots, x_{i-1}) = P(x_i|parents(X_i))$

- A conditional distribution for each node

Remark. Bayes' Net does not imply causation, only the lack thereof

8.1 Inference

Definition 8.3 (Inference by Enumeration). General procedure:

1. Join Factors: $\forall r, t : P(r, t) = P(r)P(t|r)$

2. Eliminate: $P(T) = \sum_r P(r, T)$

3. Normalize: $P(Q|e_1, \ldots, e_k) = \frac{P(Q, e_1, \ldots, e_k)}{\sum_q P(Q, e_1, \ldots, e_k)}$

Remark. Variable elimination is faster than inference by enumeration via interleaving joining and marginalizing.

8.2 Sampling

8.3 D-Separation

Definition 8.4 (Causal Chains). $X \to Y \to Z$

$$P(x, y, z) = P(x)P(y \mid x)P(z \mid y) \implies \neg(X \perp\!\!\!\perp Z) \land (X \perp\!\!\!\perp Z \mid Y)$$

Definition 8.5 (Common Causes). $X \leftarrow Y \to Z$

$$P(x, y, z,) = P(y)P(x \mid y)P(z \mid y) \implies \neg(X \perp\!\!\!\perp Z) \land (X \perp\!\!\!\perp Z \mid Y)$$

Definition 8.6 (Common Effect). $X \to Z \leftarrow Y$

$$(X \perp\!\!\!\perp Z) \land \neg(X \perp\!\!\!\perp Z \mid Y)$$

Definition 8.7 (D-Separation). $X \perp\!\!\!\perp Y \mid Z$ *iff* X and Y are "d-separated" by Z (i.e., all undirected path from X to Z are inactive). A path is active if all triple it contains is active:

Algorithm 7: Sampling.

---

1 Function PRIOR-SAMPLING:
2   **for** $i = 1, 2, \ldots, n$ *in topological order* **do**
3   Sample x_i from $P(X_i \mid \text{Parents}(X_i))$
4   **return** $(x_1, x_2, \ldots, x_n)$
5 Function REJECTION-SAMPLING:
     // Will reject lots of samples if evidence is unlikely.
6   **for** $i = 1, 2, \ldots, n$ *in topological order* **do**
7   Sample x_i from $P(X_i \mid \text{Parents}(X_i))$
8   **if** x_i *inconsistent with evidence* **then**
9   **return** // no sample generated
10   **return** $(x_1, x_2, \ldots, x_n)$
11 Function LIKELIHOOD-WEIGHTING-SAMPLING:
12 $w \leftarrow 1.0$
13   **for** $i = 1, 2, \ldots, n$ *in topological order* **do**
14   **if** X_i *is an evidence variable* **then**
15 $X_i \leftarrow$ observation x_i from X_i
16 $w \leftarrow w \times P(x_i \mid \text{Parents}(X_i))$
17   **else**
18   Sample x_i from $P(X_i \mid \text{Parents}(X_i))$
19   **return** $(x_1, x_2, \ldots, x_n), w$

- Causal chain $A \to B \to C$ where B is unobserved

- Common cause $A \leftarrow B \to C$ where B is unobserved

- Common effect $A \to B \leftarrow C$ where B or one of its descendants is observed