

Introduction to Artificial Intelligence

Daniel Deng

1 Search Problems

Definition 1.1 (Reflex Agent). A reflex agent chooses actions based on its current perception of the world.

Definition 1.2 (Planning Agent). A planning agent chooses actions based on hypothesized consequences of actions.

Definition 1.3 (Search Problem). A search problem consists of a state space, a successor function, a start state, and a goal test.

2 Search Algorithms

2.1 Heuristics

Definition 2.1 (Heuristic). A heuristic $h(n)$ is a function that estimates the distance from state n to the goal state for a particular search problem. It is often solutions of relaxed problems.

Definition 2.2 (Admissibility). A heuristic is admissible, or optimistic, if $0 \leq h(n) \leq h^*(n)$ where h^* is the true cost to goal state.

Definition 2.3 (Consistency). A heuristic is consistent if $h(n) - h(n+1) \leq c(n, n+1)$ where c is the cost between states n and $n+1$.

Remark. Consistency necessarily implies admissibility.

Table 1: Search algorithms.

| | Fringe | Complete | Optimal | Time | Space |
|------------------------------|--------------------|--------------------------|------------------------------|-------------------------|-------------------------|
| Depth-First Search | Stack | <i>iff</i> no cycle | No | $O(b^m)$ | $O(bm)$ |
| Breadth-First Search | Queue | Yes | <i>iff</i> uniform cost | $O(b^s)^1$ | $O(b^s)^1$ |
| Uniform Cost Search | PQ $(g(n))^2$ | <i>iff</i> positive cost | Yes | $O(b^{c^*/\epsilon})^3$ | $O(b^{c^*/\epsilon})^3$ |
| Greedy Search | PQ $(h(n))$ | - | No | - | - |
| A* Tree Search | PQ $(h(n) + g(n))$ | - | <i>iff</i> $h(n)$ admissible | - | - |
| A* Graph Search ⁴ | PQ $(h(n) + g(n))$ | - | <i>iff</i> $h(n)$ consistent | - | - |

¹ s = depth of solution.

² $g(n)$ = cumulative path cost.

³ c^*/ϵ = effective solution depth (c^* = cost of the cheapest solution; ϵ = minimum cost of cost-contour arcs).

⁴ Compared to tree search, graph search keeps a closed set of expanded states to check against to prevent duplicate expansions.

Remark. Implementation of search algorithms differ only in fringe strategies.

3 Constrained Satisfaction Problems

Definition 3.1 (Constrained Satisfaction Problems). Constrained Satisfaction Problems (CSPs) are a type of **identification problem** defined by variable X_0, \dots, X_n with values from a domain D that satisfies a set of constrains.

3.1 Filtering

Definition 3.2 (Arc Consistency).

$$\text{Arc } X \rightarrow Y \text{ is consistent} \Leftrightarrow (\forall x \in D_x)(\exists y \in D_y)(y \text{ can be assigned to } Y \text{ without violating a constraint.})$$

3.2 Ordering

Definition 3.3 (Minimum Remaining Values). The MRV policy chooses an unassigned variable that has the fewest valid remaining values in order to induce backtracking earlier and reduce potential node expansions.

Definition 3.4 (Least Constraining Value). The LCV policy chooses a value assignment that violates the least amount of constraints, which requires additional computation such as running arc consistency test on each value.

Algorithm 1: Backtracking search.

Input: A constraint satisfaction problem P .
Output: A complete assignment A .

```

1 Function BS( $P$ ):
2   return EXPLORE ( $P$ ,  $\{\}$ )
3 Function EXPLORE( $P$ ,  $A$ ):
4   if  $A$  is complete then
5     return  $A$ 
6    $unassigned \leftarrow$  an unassigned VARIABLES( $P$ )
7   foreach  $value \in$  DOMAIN( $unassigned$ ) do
8     if  $value$  is consistent with all CONSTRAINTS( $P$ ) then
9       add  $\{unassigned \leftarrow value\}$  to  $A$ 
10       $attempt \leftarrow$  EXPLORE( $P$ ,  $A$ )
11      if  $attempt$  failed then
12        remove  $\{unassigned \leftarrow value\}$  from  $A$ 
13      else
14        return  $attempt$ 
15   return failed

```

3.3 Structure

Given a tree-structured CSP, represent it as a directed acyclic graph. Enforcing arc consistency in reverse topological order then assigning in topological order ensures a runtime of $O(nd^2)$ (as opposed to $O(d^n)$ in the general case).

TODO: nearly tree-like CSPs and tree decomposition.

4 Local Search

Improve a single option until no further improvements can be made.

Remark. Generally, local search is faster and more memory efficient at the expense of completeness and optimality.

4.1 Iterative Algorithm for CSP/Hill Climbing

Starting with a "complete" state, randomly select any conflicted variables and reassign values using min-conflicts heuristics.

Remark. Efficiency of the algorithm depends on $R = \frac{\text{number of constraints}}{\text{number of variables}}$; computation time is approximately constant time except when R approaches the *critical ratio*.

Algorithm 2: Arc consistency filtering.

Input: A constraint satisfaction problem P .

```

1 Function AC-3( $P$ ):
2    $Q \leftarrow$  empty Queue
3   enqueue all arcs  $\in P$ 
4   while  $Q$  is not empty do
5      $(X_i, X_j) \leftarrow$  dequeue from  $Q$ 
6     if FILTER( $X_i, X_j$ ) is successful then
7       foreach  $X_k \in \text{NEIGHBORS}(X_i)$  do
8         enqueue  $(X_k, X_i)$ 
9 Function FILTER ( $tail, head$ ):
10   $result \leftarrow$  false
11  foreach  $value \in \text{DOMAIN}(tail)$  do
12    if  $value$  violates some constraint with all values in  $\text{DOMAIN}(head)$  then
13      delete  $value$  from  $\text{DOMAIN}(tail)$ 
14       $result \leftarrow$  true
15  return  $result$ 

```

4.2 Simulated Annealing

5 Genetic Algorithms

Keep best N hypotheses at each step (selection) based on a fitness function and have pairwise crossover operations (and, optionally, mutation operations) to generate a new set of hypotheses.

6 Games

Remark. Typically, performing Minimax and Expectimax to terminal states are too costly, which can be solved by terminating early and estimating the state utility with evaluation functions.

7 Markov Decision Processes

Remark. Finite horizons (finite timestep before an agent terminates) and/or discount factors (γ) ensure an agent terminates in MDP.

Definition 7.1 (Transition Function). $T(s, a, s') = P(s' \mid s, a)$

Algorithm 3: Simulated annealing.

Input: A problem P and a schedule/mapping from time to "temperature" T .

Output: A solution state.

/* Escape local maxima by allowing downhill movement based on a "temperature"-dependent probabilistic function. */

```

1  $current \leftarrow$  initial state of  $P$ 
2 for  $t \leftarrow 1$  to  $\infty$  do
3    $temp \leftarrow T[t]$ 
4   if  $temp = 0$  then
5     return  $current$ 
6   else
7      $next \leftarrow$  a randomly selected successor of  $current$ 
8      $\Delta \leftarrow \text{VALUE}[next] - \text{VALUE}[current]$ 
9     if  $\Delta > 0$  then
10       $current \leftarrow next$ 
11     else
12       $currnet \leftarrow next$  with probability  $e^{\frac{\Delta}{temp}}$ 

```

Definition 7.2 (Value Iteration). Initialize $V_0(s) \leftarrow 0$ for all $s \in S$. Compute $\forall s \in S$

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')] \quad (\text{value iteration})$$

until $V_i(s)$ converges for all $s \in S$ (yields V^*). Then, compute $\forall s \in S$

$$\pi^*(s) \leftarrow \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')] \quad (\text{policy extraction})$$

Definition 7.3 (Policy Iteration). Define an initial policy π_0 (can be arbitrary, but ideality close to the optimal policy). Then, iteratively solve $\forall s \in S$

$$V^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V^{\pi_i}(s')] \quad (\text{policy evaluation})$$

$$\pi_{i+1}(s) \leftarrow \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')] \quad (\text{policy improvement})$$

until $\pi(s)$ converges for all $s \in S$ (yields π^*).

Remark. Policy evaluation solves a system of $|S|$ linear equations.

Remark. Policy iteration converges faster than value iteration.

Algorithm 4: Minimax with alpha-beta pruning.

Input: A game state S
Output: The root minimax value.
 /* initialize $\alpha \leftarrow -\infty$ and $\beta \leftarrow \infty$ */

```

1 Function VALUE( $S, \alpha, \beta$ ):
2   if  $S$  is a terminal state then
3     return known terminal value
4   if the agent is maximizing then
5     return MAX-VALUE( $S, \alpha, \beta$ )
6   if the agent is minimizing then
7     return MIN-VALUE( $S, \alpha, \beta$ )

8 Function MAX-VALUE( $S, \alpha, \beta$ ):
9    $v \leftarrow -\infty$ 
10  foreach successor  $S'$  of  $S$  do
11     $v \leftarrow \text{MAX}(v, \text{VALUE}(S'))$ 
12    if  $v \geq \beta$  then
13      return  $v$ 
14     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
15  return  $v$ 

16 Function MIN-VALUE( $S, \alpha, \beta$ ):
17   $v \leftarrow \infty$ 
18  foreach successor  $S'$  of  $S$  do
19     $v \leftarrow \text{MIN}(v, \text{VALUE}(S'))$ 
20    if  $v \leq \alpha$  then
21      return  $v$ 
22     $\beta \leftarrow \text{MIN}(\beta, v)$ 
23  return  $v$ 

```

Algorithm 5: Expectimax.

Input: A game state S

Output: The root minimax value.

```

1 Function VALUE( $S$ ):
2   if  $S$  is a terminal state then
3     return known terminal value
4   if the agent is maximizing then
5     return MAX-VALUE( $S$ )
6   if the agent is randomizing then
7     return EXP-VALUE( $S$ )

8 Function MAX-VALUE( $S$ ):
9    $v \leftarrow -\infty$ 
10  foreach successor  $S'$  of  $S$  do
11     $v \leftarrow \text{MAX}(v, \text{VALUE}(S'))$ 
12  return  $v$ 

13 Function EXP-VALUE( $S$ ):
14   $v \leftarrow 0$ 
15  foreach successor  $S'$  of  $S$  do
16     $v \leftarrow \mathbb{E}[\text{VALUE}(S')] //$  expected value
17  return  $v$ 

```
